

## 🔗 CEBD 1261 Final Project

---

Prepared by:  
Ieva Betina  
William Acosta

## 🔗 Table of Contents

---

- [Create the cluster using Docker Swarm](#)
- [Read and save the data \(HDFS, Parquet files, MongoDB\)](#)
- [Load the data with Jupyter notebook from MongoDB](#)
- [Security precautions with environment variables](#)
- [Analyze the data with Spark Dataframes](#)
- [Spark with Google Cloud Platform](#)
- [Future improvements](#)
- [Conclusion](#)

## 🔗 Summary

---

This project will demonstrate running a Spark Cluster with Docker Swarm on a single node and also on Google Cloud Platform. As part of the demonstration we will read a [London energy csv](#) with Spark and write it to HDFS using Parquet files and then also to MongoDB. From there, we will use Jupyter notebook as a front end to analyze the data. On GCP, we will demonstrate reading the csv from a bucket and saving back to the bucket in the Parquet format. Jupyter notebook will again be used to analyze the data.

## 🔗 Create a Spark Cluster

---

### 🔗 NOTE

- make sure you have enough Docker resources. Increase the memory and CPUs.
- for this project we'll be running the swarm on just one node. We could have also run these services with Docker compose since there is only one host.

Get a swarm up and running:

```
docker swarm init
```

For multi-node clusters, you might have to use a specific hostname:

```
docker swarm init --advertise-addr $(hostname -i)
```

Create a [shared overlay network](#) named "spark-network":

```
docker network create -d overlay --attachable spark-network
```

This network allows the containers to communicate and can be added to any new containers outside of the swarm to communicate with the containers in the swarm.

## 🔗 Create a volume for files

We use a named volume "london-power" for storing files. We will be referencing this volume in our stack.yml so that the services can connect to it.

```
docker volume create london-power
```

### ALTERNATIVELY:

Don't create the volume since the yml file can create it for you at launch.

## 🔗 Create a custom docker image

We modified the mjhea0/spark:2.4.1 image to already include our python scripts. We created a simple [Dockerfile](#) with the following:

```
FROM mjhea0/spark:2.4.1
COPY script/ /usr/spark-2.4.1/script
```

Then we build the image. We added the file to docker hub so we included a user name and tags:

```
docker build -t wacosta/spark-custom:2.4.1 .
```

Push to docker hub:

```
docker push wacosta/spark-custom:2.4.1
```

The script files will now be located in the workers when the stack launches.

### Alternatives

Getting a file directly from a website either with `ADD` or `RUN wget` then unzipping them with `RUN unzip` would be a great way to automate updating of the files by using this custom Docker image when deploying our stack. Just exploring that it is possible for educational purposes. However, the documentation recommends against it and large files should be put directly in a volume :

```
FROM mjhea0/spark:2.4.1

RUN apt-get update

RUN apt-get install -y wget

RUN mkdir /mongo

USER root

RUN wget "https://data.london.gov.uk/download/smartmeter-energy-use-data-in-london-households/3527bf39-d93e-4071-8451-df2ade1ea4f2/Power-Networks-LCL-June2015.zip" -P /usr/spark-2.4.1/volume/

RUN unzip /usr/spark-2.4.1/volume/Power-Networks-LCL-June2015.zip
```

We can see that the file has unzipped and is inside the spark worker:

```
| cce-fb111-02m:2019-spring-group-project-ieva-william i_beti$ docker exec -it aa08b39161f8 sh
# ls
LICENSE NOTICE Power-Networks-LCL-June2015(withAcornOps)v2.csv R README.md RELEASE bin conf data examples jars kubernetes licenses python sbin volume work yarn
#
```

Another interesting option is to add files to the image from the cloud using AWS s3 bucket or gcp bucket. It would be worth exploring.

## 🔗 Prepare the stack yml file

The [stack-spark.yml](#) file must include all the services necessary for this project. This includes:

- spark master
- spark worker (using custom image)
- hadoop

- jupyter notebook
- mongoDB

It is also necessary to stipulate the named volume "london-power" that some of the containers will be using as well as the overlay network "spark-network".

```
volumes:
  - "london-power:/usr/spark-2.4.1/volume"
# AND add this at the end (external necessary for existing volumes)
volumes:
  london-power:
    external: true

networks:
  default:
    external:
      name: spark-network
```

## 🔗 Set your environment variables before deploying

The mongo and jupyter services make use of environment variables in the yml file. These will be discussed in the [environment variables](#) section below.

```
export ME_CONFIG_MONGODB_ADMINUSERNAME=root
export ME_CONFIG_MONGODB_ADMINPASSWORD=example
export MONGO_INITDB_ROOT_USERNAME=root
export MONGO_INITDB_ROOT_PASSWORD=example
export JUPYTER_TOKEN='project'
```

## 🔗 Deploy the stack

```
docker stack deploy --compose-file=stack-spark.yml spark
```

If necessary scale the workers up or down. The default we set is 3:

```
docker service scale spark_worker=5
```

## 🔗 Populate the volume

Use the busybox hack to copy files to a volume

```
docker run --rm -v "$(pwd)"/data:/data -v london-power:/volume busybox cp -r data/ /volume
```

If you don't write scripts directly into the image, add them this way:

```
docker run --rm -v "$(pwd)"/script:/script -v london-power:/volume busybox cp -r script/ /volume
```

## 🔗 Look inside the volume

Use the busybox hack to peak inside and see the files

```
docker run -it --rm -v london-power:/vol busybox sh
```

### NOTE:

If using the busybox hack, the volume must be populated *after* the stack is deployed. Otherwise, permission in the jupyter notebook will be denied when trying to start a new notebook.

## 🔗 Read and save the data (HDFS, Parquet files, MongoDB)

---

## 🔗 Run the .py file with spark-submit to save csv to HDFS and MongoDB

The [csv-hdfs-mongo.py](#) reads the csv as a dataframe in Spark. Then it gives it a schema and the proper headers which are specified by us. It saves the dataframe to HDFS as a parquet file. The parquet file is then loaded as a dataframe and saved to MongoDB. This is all done using spark submit. The mongo input and output uri are specified with environment variables for security.

We wrote the files to HDFS *and* MongoDB for demonstration purposes. We would use one or the other for regular use. We've also included other reference scripts for the following:

- csv to MongoDB
- csv to hdfs
- hdfs to MongoDB

Note that the spark submit needs the MongoDB connector passed as a package as well as the input/output uri passed with a --conf flag.

```
docker run -t --rm \
-v london-power:/usr/spark-2.4.1/volume \
--network=spark-network \
wacosta/spark-custom:2.4.1 \
bin/spark-submit \
--master spark://master:7077 \
--class endpoint \
--packages org.mongodb.spark:mongo-spark-connector_2.11:2.4.0 \
--conf "spark.mongodb.input.uri=mongodb://" "$MONGO_INITDB_ROOT_USERNAME:$MONGO_INITDB_ROOT_PASSWORD"@mongo/power.data?
authSource=admin" \
--conf "spark.mongodb.output.uri=mongodb://" "$MONGO_INITDB_ROOT_USERNAME:$MONGO_INITDB_ROOT_PASSWORD"@mongo/power.data?
authSource=admin" \
/usr/spark-2.4.1/script/csv-hdfs-mongo.py
```

View the results in HDFS at <http://localhost:50070>

View the results in MongoDB at <http://localhost:8181/>

HadoopOverviewDatanodesDatanode Volume FailuresSnapshotStartup ProgressUtilities

# Browse Directory

/data/power

Go!

Show25entries

Search:

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	root	supergroup	0 B	Jun 06 19:51	3	128 MB	_SUCCESS
-rw-r--r--	root	supergroup	6.71 MB	Jun 06 19:32	3	128 MB	part-00000-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	supergroup	6.63 MB	Jun 06 19:32	3	128 MB	part-00001-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	supergroup	6.57 MB	Jun 06 19:32	3	128 MB	part-00002-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	supergroup	6.52 MB	Jun 06 19:32	3	128 MB	part-00003-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	supergroup	6.64 MB	Jun 06 19:33	3	128 MB	part-00004-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	supergroup	6.77 MB	Jun 06 19:33	3	128 MB	part-00005-632feb85-401f-4c4c-baa4-b5047b06442e-c000.snappy.parquet
-rw-r--r--	root	superaroup	6.57	Jun 06 19:33	3	128 MB	part-00006-632feb85-401f-4c4c-baa4-b5047b06442e-

## Viewing Collection: data

[New Document](#)[New Index](#)

Simple

[Advanced](#)

Key

Value

String

[Find](#)[Delete all 167926914 documents retrieved](#)[← First](#)[← Prev](#)[Next →](#)[Last →](#)

_id	LCLid	stdorToU	DateTime	KWH_per_hh_per_30min	Acorn	Acorn_grouped
<a href="#">5cf9a703dfca210043613425</a>	MAC004457	Std	Mon Dec 17 2012 22:30:00 GMT+0000 (UTC)	0.393	ACORN-H	Comfortable
<a href="#">5cf9a703dfca210043613426</a>	MAC004457	Std	Mon Dec 17 2012 23:00:00 GMT+0000 (UTC)	0.386	ACORN-H	Comfortable
<a href="#">5cf9a703dfca210043613427</a>	MAC004457	Std	Mon Dec 17 2012 23:30:00 GMT+0000 (UTC)	0.423	ACORN-H	Comfortable
<a href="#">5cf9a703dfca210043613428</a>	MAC004457	Std	Tue Dec 18 2012 00:00:00 GMT+0000 (UTC)	0.296	ACORN-H	Comfortable
<a href="#">5cf9a703dfca210043613429</a>	MAC004457	Std	Tue Dec 18 2012 00:30:00 GMT+0000 (UTC)	0.131	ACORN-H	Comfortable
<a href="#">5cf9a703dfca21004361342a</a>	MAC004457	Std	Tue Dec 18 2012 01:00:00 GMT+0000 (UTC)	0.106	ACORN-H	Comfortable
<a href="#">5cf9a703dfca21004361342b</a>	MAC004457	Std	Tue Dec 18 2012 01:30:00 GMT+0000 (UTC)	0.135	ACORN-H	Comfortable
<a href="#">5cf9a703dfca21004361342c</a>	MAC004457	Std	Tue Dec 18 2012 02:00:00 GMT+0000 (UTC)	0.084	ACORN-H	Comfortable
<a href="#">5cf9a703dfca21004361342d</a>	MAC004457	Std	Tue Dec 18 2012 02:30:00 GMT+0000 (UTC)	0.036	ACORN-H	Comfortable
<a href="#">5cf9a703dfca21004361342e</a>	MAC004457	Std	Tue Dec 18 2012 03:00:00 GMT+0000 (UTC)	0.063	ACORN-H	Comfortable

1

2

3

### Remove files from HDFS

If desired, you can remove files from HDFS:

```
docker exec -it $(docker ps --filter name=master --format "{{.ID}}") hdfs dfs -rm -r hdfs://hadoop:8020/data/power/
```

For processing the full 11GB file we used 5 workers using 1 core each and 2gb of memory each. The results were much faster after scaling up for the simple process of reading, setting a schema and writing the file. Notice that before scaling up, the process took 40 minutes on 2 cores and 1gb memory each.



## Spark Master at spark://master:7077

URL: spark://master:7077

Alive Workers: 2

Cores in use: 2 Total, 0 Used

Memory in use: 2.0 GB Total, 0.0 B Used

Applications: 0 Running, 1 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

### Workers (2)

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20190606231405-10.0.0.8-45591</a>	10.0.0.8:45591	ALIVE	1 (0 Used)	1024.0 MB (0.0 B Used)
<a href="#">worker-20190606231405-10.0.0.9-38087</a>	10.0.0.9:38087	ALIVE	1 (0 Used)	1024.0 MB (0.0 B Used)

### Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

### Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
<a href="#">app-20190606233148-0000</a>	csv-hdfs-mongo	2	1024.0 MB	2019/06/06 23:31:48	root	FINISHED	40 min

After scaling up the process was much faster: 13 min using 5 cores and 2GB of memory per worker



## Spark Master at spark://master:7077

URL: spark://master:7077

Alive Workers: 5

Cores in use: 5 Total, 0 Used

Memory in use: 10.0 GB Total, 0.0 B Used

Applications: 0 Running, 1 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

### Workers (5)

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20190607130958-10.0.0.55-36811</a>	10.0.0.55:36811	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
<a href="#">worker-20190607130958-10.0.0.62-39269</a>	10.0.0.62:39269	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
<a href="#">worker-20190607131002-10.0.0.56-35305</a>	10.0.0.56:35305	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
<a href="#">worker-20190607131003-10.0.0.53-43671</a>	10.0.0.53:43671	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)
<a href="#">worker-20190607131003-10.0.0.60-35327</a>	10.0.0.60:35327	ALIVE	1 (0 Used)	2.0 GB (0.0 B Used)

### Running Applications (0)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
----------------	------	-------	---------------------	----------------	------	-------	----------

### Completed Applications (1)

Application ID	Name	Cores	Memory per Executor	Submitted Time	User	State	Duration
<a href="#">app-20190607135730-0000</a>	csv-hdfs-mongo	5	1024.0 MB	2019/06/07 13:57:30	root	FINISHED	13 min

## Load the data with Jupyter notebook from MongoDB

Use jupyter notebook to load the data from MongoDB and then run some queries.

Go to Jupyter at <http://localhost:8888>

Type in the password that was used as the environment variable, in this case "project".

Alternatively:

Grab the token by looking at the logs:

```
docker service logs spark_jupyter
```

Use the [startup\\_default\\_mongo.ipynb](#) notebook to get started. Notice that in a notebook you must specify the input and output uri and the spark connector.

```
.config('spark.jars.packages','org.mongodb.spark:mongo-spark-connector_2.11:2.4.0' ) \  
.config("spark.mongodb.input.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin") \  
.config("spark.mongodb.output.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin")
```

The screenshot shows a Jupyter Notebook interface with the title 'startup\_default\_mongo'. The top bar indicates 'Last Checkpoint: a minute ago (autosaved)' and includes a 'Logout' button. The notebook has a toolbar with icons for file operations, running, and markdown. The main content area displays four code cells:

```
In [1]: user = %env MONGO_INITDB_ROOT_USERNAME  
password = %env MONGO_INITDB_ROOT_PASSWORD  
login = user + ":" + password
```

```
In [2]: from pyspark.context import SparkContext  
from pyspark.sql.session import SparkSession  
  
spark = SparkSession \  
    .builder \  
    .appName("jupyter") \  
    .master("local") \  
    .config('spark.jars.packages','org.mongodb.spark:mongo-spark-connector_2.11:2.4.0' ) \  
    .config("spark.mongodb.input.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin") \  
    .config("spark.mongodb.output.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin") \  
    .getOrCreate()
```

```
In [3]: df = spark.read.format("com.mongodb.spark.sql.DefaultSource").load()
```

```
In [4]: df.show()
```

The output of the fourth cell shows a table with 20 rows of data. The columns are: Acorn, Acorn\_grouped, DateTime, KWH\_per\_hh\_per\_30min, LCLid, \_id, and stdorToU. The data represents energy consumption for various acorn groups over time.

Acorn	Acorn_grouped	DateTime	KWH_per_hh_per_30min	LCLid	_id	stdorToU
ACORN-U	ACORN-U	2012-05-03 05:30:00	0.114	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 06:00:00	0.177	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 06:30:00	0.165	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 07:00:00	0.307	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 07:30:00	0.547	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 08:00:00	0.497	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 08:30:00	0.657	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 09:00:00	0.16	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 09:30:00	0.107	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 10:00:00	0.107	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 10:30:00	0.106	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 11:00:00	0.106	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 11:30:00	0.168	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 12:00:00	0.456	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 12:30:00	0.442	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 13:00:00	0.659	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 13:30:00	0.62	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 14:00:00	0.523	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 14:30:00	1.15	MAC000023	[5cf85c0a29582e00...	Std
ACORN-U	ACORN-U	2012-05-03 15:00:00	0.905	MAC000023	[5cf85c0a29582e00...	Std

only showing top 20 rows

## Save results from Jupyter to MongoDB

### Save to default output uri database and collection

```
df.write.format("com.mongodb.spark.sql.DefaultSource").mode("append").save()
```

### Save to a different db and coll

You must specify the output uri in the config before doing this. For a db:results and coll:data

```
.config("spark.mongodb.output.uri", "mongodb://" + login + "@mongo/results.data?authSource=admin")
```

Then write to that destination:

```
df.write.format("com.mongodb.spark.sql.DefaultSource").mode("append").option("database",
    "results").option("collection", "data").save()
```

Mongo Express Database: power Collection: results

Viewing Collection: results

New DocumentNew Index

SimpleAdvanced

Key

Value

String

Find

Delete all 56 documents retrieved

First

PrevNext

Last

_id	Month	stdorToU	avg_KWH
5cfaf9c4f1eb3100c4db98f0	2014-2	Std	0.23669655507468515
5cfaf9c4f1eb3100c4db98f1	2014-2	ToU	0.21835231193535148
5cfaf9c420faf600c4ebc04a	2014-1	ToU	0.22395989790339255
5cfaf9c420faf600c4ebc04b	2014-1	Std	0.24605654569799618
5cfaf9c5f8d25700c6afb189	2013-9	ToU	0.17069914215159657
5cfaf9c5f8d25700c6afb18a	2013-9	Std	0.1856881813159775
5cfaf9c4bb6d1300c7d0ad37	2012-4	ToU	0.2076621374614184
5cfaf9c4bb6d1300c7d0ad38	2012-4	Std	0.22002929090622506
5cfaf9c5f8d25700c6afb18b	2013-6	Std	0.1768745225823871
5cfaf9c5f8d25700c6afb18c	2013-6	ToU	0.16286682419652868

### Security precautions with environment variables

Set environment variables in your local shell before deploying the stack.

```
export ME_CONFIG_MONGODB_ADMINUSERNAME=root
export ME_CONFIG_MONGODB_ADMINPASSWORD=example
export MONGO_INITDB_ROOT_USERNAME=root
export MONGO_INITDB_ROOT_PASSWORD=example
export JUPYTER_TOKEN='project'
```

Include these in the stack yml file. The values will be added to each service where they are specified

```
jupyter:
  image: jupyter/pyspark-notebook:f646d2b2a3af
```



```

volumes:
  - london-power:/home/jovyan/work
ports:
  - 8888:8888
environment:
  - MONGO_INITDB_ROOT_USERNAME
  - MONGO_INITDB_ROOT_PASSWORD
  - JUPYTER_TOKEN

mongo:
  image: mongo
  restart: always
  environment:
    - MONGO_INITDB_ROOT_USERNAME
    - MONGO_INITDB_ROOT_PASSWORD

mongo-express:
  image: mongo-express
  restart: always
  ports:
    - 8181:8081
  environment:
    - ME_CONFIG_MONGODB_ADMINUSERNAME
    - ME_CONFIG_MONGODB_ADMINPASSWORD

```

We used them to run the spark-submit in order to specify the uri. This allowed us to avoid using credentials in the py file itself

```

--conf "spark.mongodb.input.uri=mongodb://" + MONGO_INITDB_ROOT_USERNAME + ":" + MONGO_INITDB_ROOT_PASSWORD + "@mongo/power.data?authSource=admin" \
--conf "spark.mongodb.output.uri=mongodb://" + MONGO_INITDB_ROOT_USERNAME + ":" + MONGO_INITDB_ROOT_PASSWORD + "@mongo/power.data?authSource=admin" \
/usr/spark-2.4.1/script/csv-hdfs-mongo.py

```

We then also use them in the jupyter notebook to avoid obviously exposing the credentials. They can still be accessed but they're not written directly into the notebook at least.

```

user = %env MONGO_INITDB_ROOT_USERNAME
password = %env MONGO_INITDB_ROOT_PASSWORD
login = user + ":" + password

```

```

.config("spark.mongodb.input.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin") \
.config("spark.mongodb.output.uri", "mongodb://" + login + "@mongo/power.data?authSource=admin") \

```

An alternative for the jupyter file would have been to set the uri in the spark stack so it wouldn't have to be referenced in a cell. This is like the spark-submit approach. Unfortunately, we weren't sure how to set the uri at runtime. We also would like to further investigate Docker Secrets.

The os.environ.get is also useful within python files. The magic %env that we used is better for notebooks.

```

import os
user = os.environ.get('MONGO_USER')
password = os.environ.get('MONGO_PASS')
login = user + ":" + password

```

## 🔗 Analyze the data with Spark Dataframes

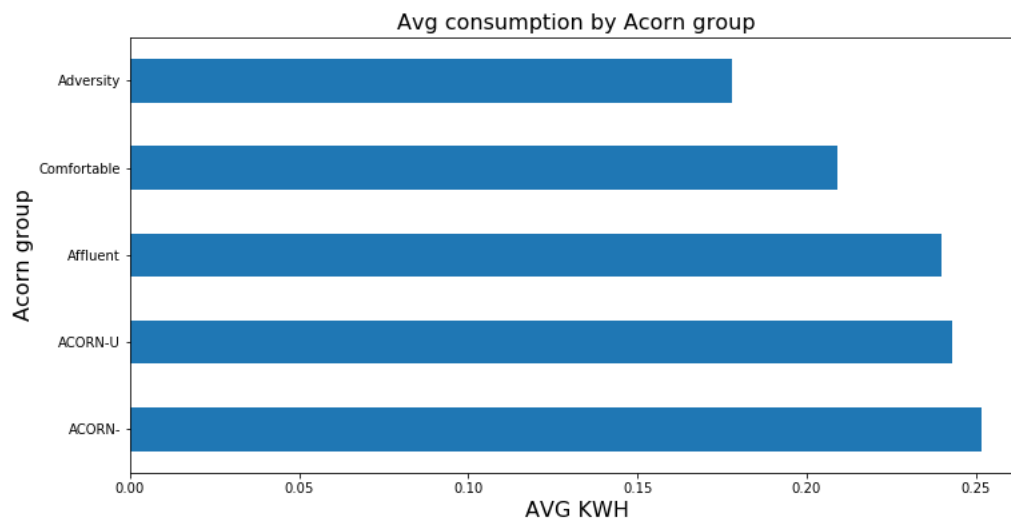
We used Spark to connect to our MongoDB with all the data that we uploaded. For the full 11GB file however, we found it faster to read from HDFS so we used MongoDB to save the results instead. We used both pythonic syntax and SparkSQL syntax to demonstrate.

See [analysis\\_final.ipynb](#) for the specifics.


In summary, we read the file in from HDFS as a dataframe:

```
df = spark.read.parquet('hdfs://hadoop:8020/data/power')
```

We confirmed there were 167,926,914 records. We then set out to find the monthly consumption patterns of each Acorn group and plot them. On average, the energy consumption was highest in "affluent" households, followed by "comfortable" and lowest in "adversity" households.



Furthermore, we wanted to see if the test group that received warnings about energy prices beforehand consumed less energy than those who did not. The results showed that they did.



Finally, we saved the results to a new MongoDB collection:

```
results.write.format("com.mongodb.spark.sql.DefaultSource").mode("append").option("database", \
"power").option("collection", "results").save()
```

## Spark with Google Cloud Platform

### Running a spark submit on GCP

Create project called cebd1261

Create a cluster called project with 2 workers

Search clusters, press Enter

<input type="checkbox"/>	Name ^	Region	Zone	Total worker nodes	Scheduled deletion	Cloud Storage staging bucket	Created	Status
<input type="checkbox"/>	<input checked="" type="checkbox"/> project	global	us-east1-b	2	Off	cebd1261	Jun 5, 2019, 7:39:55 PM	Running

Create and populate the bucket with /data and /script using Regional Storage as Storage Class .This class is optimal for our need as it is for storing frequently accessed data in the same region as Google CloudProc Or Google Compute Engine instances that use it such as data analytics and is lower cost per GB than Multi-Regional Storage used in class tutorial.

We add a label to our bucket for example 'project' so if we theoretically have many buckets they all are labeled for better access.

We also can Edit Permissions to our bucket to add another user (like the other teammate for the project ) so they can see what is inside inside the bucket (it can be read-only permission as well)

To work with storage through the command line we use gsutil

```
- `gsutil ls gs://cebd1261/` to see what is inside the bucket
- `gsutil ls -L` gives a lot information about the settings of the bucket
- `gsutil mb gs://cebd1261_project/` to create a new bucket
- `gsutil mv gs://cebd1261/gcp.py gs://cebd1261_project/` to move a file from one bucket to another
- ...
```

Storage

Browser

Transfer

Transfer Appliance

Settings

Bucket details

EDIT BUCKET

REFRESH BUCKET

cebd1261

Objects

Overview

Permissions

Bucket Lock

Upload files

Upload folder

Create folder

Manage holds

Delete

Filter by prefix...

Buckets / cebd1261

<input type="checkbox"/>	Name	Size	Type	Storage class	Last modified	Public access	Encryption	Retention expiration date	Holds
<input type="checkbox"/>	.ipynb_checkpoints/	—	Folder	—	—	Per object	—	—	-
<input type="checkbox"/>	data/	—	Folder	—	—	Per object	—	—	-
<input type="checkbox"/>	google-cloud-dataproc-metainfo/	—	Folder	—	—	Per object	—	—	-
<input type="checkbox"/>	notebooks/	—	Folder	—	—	Per object	—	—	-
<input type="checkbox"/>	script/	—	Folder	—	—	Per object	—	—	-

Create and submit a spark job by choosing PySpark as Job type and pointing to our bucket for PyScript to run

Dataproc

Submit a job

Clusters

Jobs

Workflows

Job ID

job-e51035d3

Region

global

Cluster

project

Job type

PySpark

Main python file

gs://cebd1261/script/gcp.py




Additional python files (Optional)

Enter file path, for example, hdfs://example/example.py


Arguments (Optional)

1000

See Also (Optional)

	Search jobs, press Enter						
<input type="checkbox"/> Job ID	Region	Type	Cluster	Start time	Elapsed time	Status	
<input type="checkbox"/>  job-48f03a00	global	PySpark	jevfka-proj	Jun 5, 2019, 9:17:49 PM	28 sec	Succeeded	

When Job output is complete we can see the results


**job-48f03a00**

Start time: Jun 5, 2019, 9:17:49 PM
 Elapsed time: 28 sec
 Status:

Output
 Configuration

☐ Line wrapping
 Equivalent command

```
19/06/06 01:18:00 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1559783738287_0001
```

	LCliId stdo TolU	DateTime KWH_per_hh_per_30min	Acorn Acorn_grouped
[MAC000002]	Std 2012-10-12 00:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 01:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 01:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 02:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 02:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 03:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 03:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 04:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 04:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 05:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 05:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 06:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 06:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 07:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 07:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 08:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 08:30:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 09:00:00	0.0 ACORN-A	Affluent
[MAC000002]	Std 2012-10-12 09:30:00	0.0 ACORN-A	Affluent

## BigQuery

BigQuery on GCP can process petabytes in minutes. To query data we pin the project associated with data we are using and create a table. Then populate the table from the Google Storage bucket where our data is stored:

## Create table

### Source

Create table from: Google Cloud Storage Select file from GCS bucket: cebd1261/power.csv File format: CSV

### Destination

Project name: CEBD1261 Dataset name: power\_csv Table type: Native table

### Table name

group\_project\_2019

### Schema

#### Auto detect

☒ Schema and input parameters

Schema will be automatically generated.

### Partition and cluster settings

#### Partitioning:

No partitioning

#### Clustering order (optional):

Clustering order determines the sort order of the data. Clustering can only be used on a partitioned table, and works with tables partitioned either by column or ingestion time.

Comma-separated list of fields to define clustering order (up to 4)

It is possible to load data from Google Cloud bucket as well as from a local machine and other sources. Adding data is incredibly easy. We had some trouble with `nulls`. To solve this in `Advanced settings` in table creating pane click `ignore unrecognized characters` the database was created and nulls ignored.

The screenshot displays the Google Cloud Platform BigQuery interface. On the left, the 'Resources' panel shows the project 'project' and the dataset 'power\_csv'. The main area shows the 'Query editor' with a table preview. The table has 22 rows and 7 columns: Row, LCLid, stdorTou, DateTime, KWH\_hh\_per\_half\_hour, Acorn, and Acorn\_grouped. The data shows a sequence of rows with varying KWH values and Acorn groupings.

Row	LCLid	stdorTou	DateTime	KWH_hh_per_half_hour	Acorn	Acorn_grouped
1	MAC000200	Std	00:00.0	0.0	ACORN-F	Comfortable
2	MAC000200	Std	00:00.0	0.0	ACORN-F	Comfortable
3	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
4	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
5	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
6	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
7	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
8	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
9	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
10	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
11	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
12	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
13	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
14	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
15	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
16	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
17	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
18	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
19	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
20	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
21	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable
22	MAC000200	Std	00:00.0	0.125	ACORN-F	Comfortable

Running a query using a simple SQL language. The queries can be scheduled to run as needed. There are two types of pricing for running queries either flat rate or per query. After running each query the pane shows how many Kb the process took and that way it is easy to plan what pricing is better.

The screenshot shows the Google Cloud Platform BigQuery interface. On the left is a sidebar with navigation options like Query history, Saved queries, Job history, Transfers, Scheduled queries, BI Engine, and Resources. The main area is the Query editor, which contains a SQL query: `SELECT Acorn_grouped, SUM(KWH_hh_per_half_hour___) as TotalKWH FROM power_csv; ORDER BY Acorn_grouped; ORDER BY TotalKWH DESC`. Below the editor, there's a 'Run' button and a status bar indicating the query is valid and will process 242.6 KB. The 'Query results' section shows a table with columns 'Row', 'Acorn\_grouped', and 'TotalKWH'. The first row shows 'Comfortable' with a value of 1060.936996999753.

## Jupyter notebooks in a Cloud Datalab Session

Cloud Datalab runs on Compute Engine and connects to many cloud services running on a vm instance. It has built-in Jupyter Notebooks that supports Tensorflow and SkLearn and can analyze terabytes of data hence often used for Machine learning.

First open secure shell and run `gcloud components install datalab` > then run `datalab create datalabvm` > choose zone (say us-east-b) > console will output the localhost port that datalab is available at > choose web preview and change port accordingly if necessary

We run a simple Machine Learning task predicting Acorn\_grouped depending on KWH per half an hour

The screenshot shows a Jupyter Notebook in Google Cloud Datalab. The code cell contains imports for pandas, numpy, matplotlib, seaborn, and sklearn, followed by data loading and preprocessing steps. The data is loaded from 'power.csv' and filtered for non-null values. The resulting DataFrame is shown below the code cell.

	LCLId	stdorTou	DateTime	KWH/hh (per half hour)	Acorn	Acorn_grouped
0	MAC000002	Std	2012-10-12 00:30:00.0000000	0	ACORN-A	Affluent
1	MAC000002	Std	2012-10-12 01:00:00.0000000	0	ACORN-A	Affluent
2	MAC000002	Std	2012-10-12 01:30:00.0000000	0	ACORN-A	Affluent
3	MAC000002	Std	2012-10-12 02:00:00.0000000	0	ACORN-A	Affluent
4	MAC000002	Std	2012-10-12 02:30:00.0000000	0	ACORN-A	Affluent

We get a prediction Mean squared error of 0.39

The screenshot shows the continuation of the Jupyter Notebook. It displays the sklearn LinearRegression model training process and the calculation of the Mean Absolute Error (MAE). The final output shows the MAE is 0.39.

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

mae = mean_absolute_error(y_test, y_pred)
print('MAE', round(mae, 2))
```

MAE 0.39

## Install and run a Jupyter notebook on a Cloud Dataproc cluster

Create a cluster and install the Jupyter component through the shell

```
gcloud beta dataproc clusters create jupyter-cluster \
  --optional-components=ANACONDA,JUPYTER \
  --image-version=1.3 \
  --enable-component-gateway \
  --bucket cebd1261 \
  --project cebd1261-242823
```

Navigate to the [Cloud Dataproc Clusters](#), then select our cluster to open the [Cluster details](#) form. Click the [Web Interfaces](#)

Clusters

CREATE CLUSTER


REFRESH

DELETE

REGIONS

SHOW INFO PANEL

Search clusters, press Enter

<input type="checkbox"/> Name ^	Region	Zone	Total worker nodes	Scheduled deletion	Cloud Storage staging bucket	Created	Status
<input type="checkbox"/>  jupyter-cluster	global	us-east1-b	2	Off	cebd1261	Jun 6, 2019, 3:45:18 PM	Running

We can open a PySpark notebook stored in a bucket and work with data also stored inside our bucket by pointing to it during read:

```
spark.read.csv('gs://cebd1261/data/power.csv', schema=final_struc, header=True)
```

```
In [2]: from pyspark.context import SparkContext
        from pyspark.sql.session import SparkSession
        # sc = SparkContext('local')
        # spark = SparkSession(sc)

In [3]: from pyspark.sql.types import (StructField, StringType,
        IntegerType, TimestampType,
        FloatType, StructType)

In [4]: data_schema = [StructField('LCLid', StringType(), True),
        StructField('stdorToU', StringType(), True),
        StructField('DateTime', TimestampType(), True),
        StructField('KWH_per_hh_per_30min', FloatType(), True),
        StructField('Acorn', StringType(), True),
        StructField('Acorn_grouped', StringType(), True)]

        final_struc = StructType(fields=data_schema)

In [5]: power = spark.read.csv('gs://cebd1261/data/power.csv', schema=final_struc, header=True)

In [6]: power.printSchema()

root
 |-- LCLid: string (nullable = true)
 |-- stdorToU: string (nullable = true)
 |-- DateTime: timestamp (nullable = true)
 |-- KWH_per_hh_per_30min: float (nullable = true)
 |-- Acorn: string (nullable = true)
 |-- Acorn_grouped: string (nullable = true)

In [7]: power.count()

Out[7]: 1000000

In [8]: print('The total number of rows is:', power.count(), '\nThe total number of columns is:', len(power.columns))

('The total number of rows is:', 1000000, '\nThe total number of columns is:', 6)

In [9]: power.show(50)
```

LCLid	stdorToU	DateTime	KWH_per_hh_per_30min	Acorn	Acorn_grouped
MAC000002	Std	2012-10-12 00:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 01:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 01:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 02:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 02:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 03:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 03:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 04:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 04:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 05:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 05:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 06:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 06:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 07:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 07:30:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 08:00:00	0.0	ACORN-A	Affluent
MAC000002	Std	2012-10-12 08:30:00	0.0	ACORN-A	Affluent

In this Diploma we had opportunity to explore basics of SkLearn however we took this opportunity to learn a little bit about MLlib. During GCP part of this project we applied a quick Linear Regression model in sklearn to our data in DataLab trying to predict KHW/hh used based on household code and the Acorn\_grouped column that measures social status of the household. Ideally we wanted to use also the column where it indicates wheather the people were advised about the peaks in electricity beforehand, time of the day etc however the sample dataset had only one category and hence useless as a feature. When we tried to upload the 11G dataset to Hadoop the spark-submit would interrupt and the WARN would print that not enough resources were allocated (even after changing amount of spark\_workers to 3 and memory to 2g in spark-submit .yaml file ) We also tried DecisionTreeRegressor, RandomForestRegressor, GBTRRegressor with poor results.

Mean Squared error with SkLearn was 0.27 and with MLlib 0.20 which surprised us since we expected same results for both. In the future more models should be trained using more accurate feature engineering. Notebooks available in our repo.

## 🔗 Future Improvements

---

- Make multi-node cluster
- Fully automated with a scheduler like airflow
- Machine Learning with Spark MLlib
- Improve security
- Connect MongoDB to Dataproc

## 🔗 Conclusion

---

Through the project we refined our docker skills and explored various file types and storage options. We got to really see the power of distributed Spark jobs by running our script on many workers with an 11GB file. We also explored new territory on Google Cloud Platform, scratching the surface of running Spark in a cloud environment. We briefly explored the data with Spark dataframes. We look forward to using more advance Spark features like MLlib.